



Configure a DAG with Celery Executor, PostgreSQL and RabbitMQ

Time to practice!



Instructions

- If your airflow webserver and airflow scheduler are running, stop them by typing `ctrl-C` in their respective terminal.
- `vim ~/airflow/airflow.cfg`
- In the configuration change the following:
 - `executor = CeleryExecutor`
 - `sql_alchemy_conn = postgresql+psycopg2://airflow@localhost:5432/airflow_mdb`
 - `broker_url = pyamqp://admin:rabbitmq@localhost/`
 - `result_backend = db+postgresql://airflow@localhost:5432/airflow_mdb`
 - `worker_log_server_port = 8794`
- Restart airflow
 - `airflow initdb`
 - `airflow webserver`
 - `airflow scheduler`



Important Note

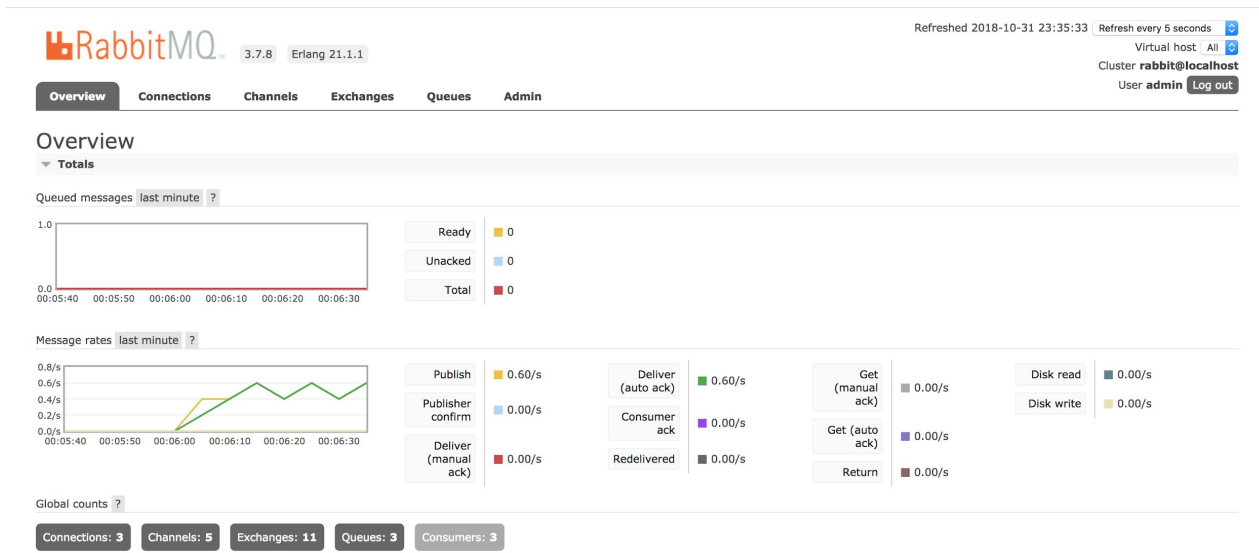
- When we type `airflow worker` we actually run a Celery Worker Node.
- When a celery worker is running, it creates one parent process to manage the running tasks. This process handles features like sending/receiving queue messages, tracking status, registering and killing tasks, etc.
- This process then spawns N number of child worker processes that actually execute the individual tasks.
- The number of child worker processes can be determined by typing: `airflow worker -c 2` (for 2 child worker processes) or by changing the `worker_concurrency` parameter in `airflow.cfg`



Instructions

- When we start an airflow worker, airflow starts a tiny web server subprocess to serve the workers local log files to airflow main web server. This tiny web server use the parameter `worker_log_server_port`
- Now your Airflow worker is started, let's take a look at the RabbitMQ interface.

RabbitMQ UI





Instructions

- Now we are going to actually reuse the DAG 'dynamic_dag' but now with Celery Executors as we have configured into airflow.cfg
- Go on the Airflow UI and turn ON the dynamic_dag's toggle. Then click on the 'trigger run' icon on the right into the box links.
- After refreshing a couple of times the Airflow UI you should see that your DAG Run ended well. So how can I see if I'm actually running my DAG in a distributed mode ?



Instructions

- There are three spots to watch, the first one is actually the worker logs:

```
[2018-10-24 23:44:22,479] {settings.py:174} INFO - setting.configure_orm(): Using pool settings. pool_size=5, pool_recycle=1800
[2018-10-24 23:44:22,499] {settings.py:174} INFO - setting.configure_orm(): Using pool settings. pool_size=5, pool_recycle=1800
[2018-10-24 23:44:22,514] {settings.py:174} INFO - setting.configure_orm(): Using pool settings. pool_size=5, pool_recycle=1800
[2018-10-24 23:44:23,894] {__init__.py:51} INFO - Using executor CeleryExecutor
[2018-10-24 23:44:23,937] {__init__.py:51} INFO - Using executor CeleryExecutor
[2018-10-24 23:44:23,919] {__init__.py:51} INFO - Using executor CeleryExecutor
[2018-10-24 23:44:24,307] {models.py:258} INFO - Filling up the DagBag from /home/airflow/airflow/dags/dynamic_dag.py
[2018-10-24 23:44:24,319] {models.py:258} INFO - Filling up the DagBag from /home/airflow/airflow/dags/dynamic_dag.py
[2018-10-24 23:44:24,328] {models.py:258} INFO - Filling up the DagBag from /home/airflow/airflow/dags/dynamic_dag.py
[2018-10-24 23:44:24,461] {cli.py:492} INFO - Running <TaskInstance: dynamic_dag.opr_insert_2 2018-10-24T23:44:19.394240+00:00 [queued]> on host localhost.localdomain
[2018-10-24 23:44:24,477] {cli.py:492} INFO - Running <TaskInstance: dynamic_dag.opr_insert_3 2018-10-24T23:44:19.394240+00:00 [queued]> on host localhost.localdomain
[2018-10-24 23:44:24,497] {cli.py:492} INFO - Running <TaskInstance: dynamic_dag.opr_insert_1 2018-10-24T23:44:19.394240+00:00 [queued]> on host localhost.localdomain
[2018-10-24 23:44:31,750] {settings.py:174} INFO - setting.configure_orm(): Using pool settings. pool_size=5, pool_recycle=1800
[2018-10-24 23:44:32,153] {__init__.py:51} INFO - Using executor CeleryExecutor
[2018-10-24 23:44:32,250] {models.py:258} INFO - Filling up the DagBag from /home/airflow/airflow/dags/dynamic_dag.py
[2018-10-24 23:44:32,280] {cli.py:492} INFO - Running <TaskInstance: dynamic_dag.opr_end 2018-10-24T23:44:19.394240+00:00 [queued]> on host localhost.localdomain
```



Instructions

- The second spot is into the PostgreSQL database:

```
airflow_mdb=# SELECT * FROM celery_taskmeta ;
```

id	task_id	status	result	date_done	traceback
20	e1d126bf-ef0e-48d4-bdb2-e8281344ec2c	SUCCESS		2018-10-25 00:03:02.640159	
21	fecf7791-01c5-4419-b15f-b3fb52356f18	SUCCESS		2018-10-25 00:03:02.701409	
22	12cdbbd8-7404-43e7-9c3b-c817093f24fc	SUCCESS		2018-10-25 00:03:02.767025	
23	66fb230c-e593-41f4-bda7-d348069ae878	SUCCESS		2018-10-25 00:03:08.788044	

```
(4 rows)
```


Instructions

- And the third spot is from the RabbitMQ UI:

